HW 6

Created @October 12, 2025 11:44 PM

Part2

Load Testing & Analysis

search service

```
# 1. run the app
cd ~/CS6650/hw6/src
python main.py

# 2. new terminal to test
curl "http://localhost:8080/products/search?q=alpha"
curl "http://localhost:8080/health"
curl "http://localhost:8080/"

# 3. docker check
docker build -t hw6-search .
# run
docker run -p 8080:8080 hw6-search
# Test
curl http://localhost:8080/health

# 4. run locust
locust -f locustfile.py --host http://localhost:8080

# 5. aws config
[default]
aws_access_key_id=ASIAVRUVWC52BNAKBXTT
aws_secret_access_key=1VE6dPuWew+VuDQS4XI4y+cDnf1mXVfaf56xjsBQ
aws_session_token=IQoJb3JpZ2luX2VjEJj//////////wEaCXVzLXdlc3QtMiJHMEUCIDaxfUmsRmW1BFgIpYgW
P4Upd69X9d7hmYiFoAVrtFCZAiEAiNtRnKVSGJCmAx91GR0fkt4nmDXEQwipsXiKicXF+KMqtQIIQRAAGgwz
ODE0OTIyNzA5NjQiDIs/aa6zCC3Gl7PkVyqSAiPqGxcgNgVtxCgJaGX/Xp2dANZyL3DOsaGArr4nhGWi7DxU
LLmDu4haOVzioa6IDc/2MhpfD+UwWUWNuuJ8fCWjHRubfnuShoBTAHilSqCNe/2b6poF2+vpzb2XUQMuf2
lt7DMELeu+Cv5LLgbKAXUC6pnoeFLdf+6WDWFsP/uKvSTY9K9hjT/pGn55pLyw8bGf6hBx4AstJf7nAKk+pZ
CSzWrpNxOgFTjEvFvKLMLVFvuD6x/U2myDbp0URymTEt+b1vqlAtNiW+/UcJAdBPNYR6pWSR5VnOkHptRf
wrfHhFHh7KGpvsm9vAl4kZB3MlqWYLOFZ5p6dVD75RQ/laO5lLBB/xXY5jRJyb7EcrIBxjMwzeWyxwY6nQEr
8XcbHIolvUUh28nptSOh/RiU/jCpBN7q4Tk4ohoL2YJDDfTh52KGVDIFn+aPFSTOvznG85SQSxW2FkM+zYtel
pE5FtgIcW6dJg8AK8dJwGbjKOEd3AlG4z0Zn4QZ55EWN2xyN+5pNPrl5/XwLRPbOffWu0nrE0WzVX1V7G5
Of6klTZSccy3SCoCr+yKAtUVdrGSw4gDeq5+gHGJd

aws configure set aws_session_token
```

```
# 验证
aws sts get-caller-identity

# 6. build ecr
aws ecr create-repository --repository-name hw6-product-search --region us-west-2

# 1) 登录到ECR
aws ecr get-login-password --region us-west-2 | docker login --username AWS --password-stdin 381492270964.dkr.ecr.us-west-2.amazonaws.com

# 2) 构建镜像（为Linux/AMD64架构）
docker build --platform linux/amd64 -t hw6-search .

# 3) 标记镜像
docker tag hw6-search:latest 381492270964.dkr.ecr.us-west-2.amazonaws.com/hw6-product-search:latest

# 4) 推送镜像
docker push 381492270964.dkr.ecr.us-west-2.amazonaws.com/hw6-product-search:latest

# 5) Terroform
terraform apply -auto-approve

# 6) IPVS and run locust
# 获取公网ip
aws ecs describe-tasks --cluster hw6-part2-cluster --tasks arn:aws:ecs:us-west-2:381492270964:task/hw6-part2-cluster/bc7df44f276d4989998d9eed5bb124d2 --region us-west-2 | jq -r '.tasks[0].containers[0].networkInterfaces[0].privateIpv4Address as $private | .tasks[0].attachments[0].details[] | select(.name=="networkInterfaceId").value' | xargs -I {} aws ec2 describe-network-interfaces --network-interface-ids {} --region us-west-2 | jq -r '.NetworkInterfaces[0].Association.PublicIp'

# test
18.237.97.111
curl -s http://18.237.97.111/health
curl -s "http://18.237.97.111:8080/products/search?q=Alpha" | jq '.'

locust -f locustfile.py --host http://18.237.97.111:8080

7) destroy terroform
terraform destroy -auto-approve
```

```
$ cd /Users/ronghuang/MyCScode/NEU/CS6650/hw6/src && python
main.py
INFO:     Started server process [45765]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://0.0.0.0:8080 (Press CTRL
+C to quit)
INFO:     127.0.0.1:53969 - "GET /health HTTP/1.1" 200 OK
INFO:     127.0.0.1:53975 - "GET /products/search?q=alpha HT
TP/1.1" 200 OK
INFO:service:Search 'alpha': checked=100, found=20, returned
=20, time=0.000s
INFO:service:Search 'alpha': checked=100, found=20, returned
=20, time=0.001s
INFO:     127.0.0.1:54066 - "GET /products/search?q=alpha HT
TP/1.1" 200 OK
INFO:service:Search 'alpha': checked=100, found=20, returned
=20, time=0.002s
INFO:     127.0.0.1:54089 - "GET /products/search?q=alpha HT
TP/1.1" 200 OK
INFO:     127.0.0.1:54089 - "GET /favicon.ico HTTP/1.1" 404
Not Found
INFO:     127.0.0.1:54113 - "GET /health HTTP/1.1" 200 OK
INFO:     127.0.0.1:54123 - "GET / HTTP/1.1" 200 OK
INFO:     127.0.0.1:54132 - "GET / HTTP/1.1" 200 OK
INFO:     127.0.0.1:54157 - "GET /health HTTP/1.1" 200 OK
```


```
t description 36","brand":"Alpha"},{"id":41,"name":"Product
Alpha 41","category":"Electronics","description":"Product de
scription 41","brand":"Alpha"},{"id":46,"name":"Product Alph
a 46","category":"Electronics","description":"Product descri
ption 46","brand":"Alpha"},{"id":51,"name":"Product Alpha 51
","category":"Electronics","description":"Product descriptio
n 51","brand":"Alpha"},{"id":56,"name":"Product Alpha 56","c
ategory":"Electronics","description":"Product description 56
","brand":"Alpha"},{"id":61,"name":"Product Alpha 61","categ
ory":"Electronics","description":"Product description 61","b
rand":"Alpha"},{"id":66,"name":"Product Alpha 66","category"
:"Electronics","description":"Product description 66","brand
":"Alpha"},{"id":71,"name":"Product Alpha 71","category":"El
ectronics","description":"Product description 71","brand":"A
lpha"},{"id":76,"name":"Product Alpha 76","category":"Electr
onics","description":"Product description 76","brand":"Alpha
"},{"id":81,"name":"Product Alpha 81","category":"Electronic
s","description":"Product description 81","brand":"Alpha"},{
"id":86,"name":"Product Alpha 86","category":"Electronics","
description":"Product description 86","brand":"Alpha"},{"id"
:91,"name":"Product Alpha 91","category":"Electronics","desc
ription":"Product description 91","brand":"Alpha"},{"id":96,
"name":"Product Alpha 96","category":"Electronics","descript
ion":"Product description 96","brand":"Alpha"}],"total_found
```


```
(base) ronghuang@Reginas-macbook CS6650 % curl "http://local
host:8080/health"
{"status":"healthy","products_count":100000}%
(base) ronghuang@Reginas-macbook CS6650 % curl "http://local
host:8080/"
{"service":"HW6 Product Search API","endpoints":["/products/
search","/health"]}%
(base) ronghuang@Reginas-macbook CS6650 %
```

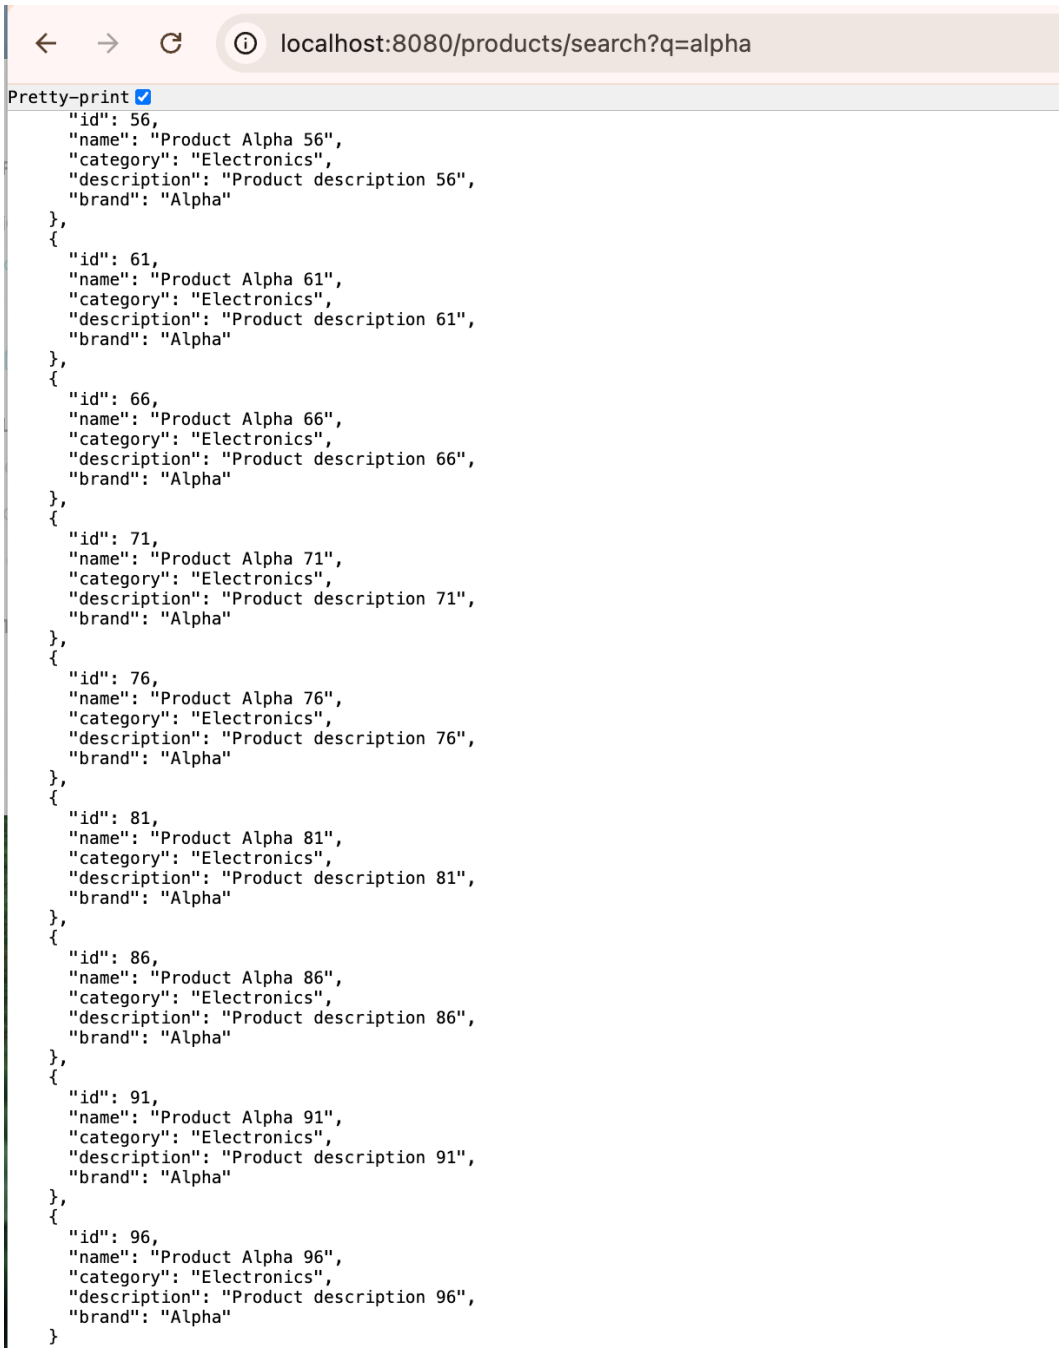
localhost:8080/products/search?q=alpha
Pretty-print
"id": 56,
"name": "Product Alpha 56",
"category": "Electronics",
"description": "Product description 56",
"brand": "Alpha"
},
{
"id": 61,
"name": "Product Alpha 61",
"category": "Electronics",
"description": "Product description 61",
"brand": "Alpha"
},
{
"id": 66,
"name": "Product Alpha 66",
"category": "Electronics",
"description": "Product description 66",
"brand": "Alpha"
},
{
"id": 71,
"name": "Product Alpha 71",
"category": "Electronics",
"description": "Product description 71",
"brand": "Alpha"
},
{
"id": 76,
"name": "Product Alpha 76",
"category": "Electronics",
"description": "Product description 76",
"brand": "Alpha"
},
{
"id": 81,
"name": "Product Alpha 81",
"category": "Electronics",
"description": "Product description 81",
"brand": "Alpha"
},
{
"id": 86,
"name": "Product Alpha 86",
"category": "Electronics",
"description": "Product description 86",
"brand": "Alpha"
},
{
"id": 91,
"name": "Product Alpha 91",
"category": "Electronics",
"description": "Product description 91",
"brand": "Alpha"
},
{
"id": 96,
"name": "Product Alpha 96",
"category": "Electronics",
"description": "Product description 96",
"brand": "Alpha"
}

docker run

```
16589a2af286b  0.0s
 => => exporting manifest list sha256:7403283143
477b7ecd35eb9  0.0s
 => => naming to docker.io/library/hw6-search:la
test           0.0s
 => => unpacking to docker.io/library/hw6-search
:latest        0.3s

 1 warning found (use docker --debug to expand):
 - FromPlatformFlagConstDisallowed: FROM --platf
orm flag should not use constant value "linux/am
d64" (line 2)

What's next:
    View a summary of image vulnerabilities and
recommendations → docker scout quickview
(base) ronghuang@Reginas-macbook hw6 % docker ru
n -p 8080:8080 hw6-search
WARNING: The requested image's platform (linux/a
md64) does not match the detected host platform
(linux/arm64/v8) and no specific platform was re
quested
INFO:     Started server process [1]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://0.0.0.0:8080
 (Press CTRL+C to quit)
INFO:     172.17.0.1:64358 - "GET /health HTTP/1
.1" 200 OK
```

```
  pyenv shell 3.8.16
  (base) ronghuang@Reginas-macbook hw6 % pyenv she
ll 3.8.16
  pyenv: shell integration not enabled. Run `pyenv
   init' for instructions.
  (base) ronghuang@Reginas-macbook hw6 % curl http
://localhost:8080/health
  {"status":"healthy","products_count":100000
(base) ronghuang@Reginas-macbook hw6 %
```

aws running

54.203.20.166

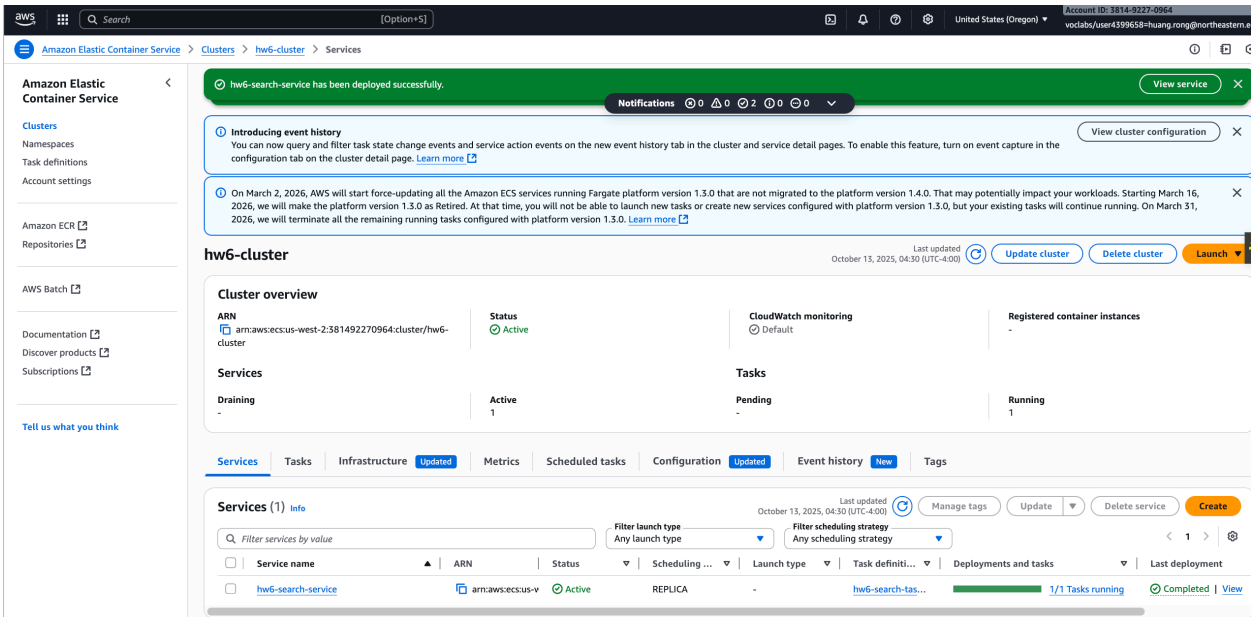

```
(base) ronghuang@Reginas-macbook CS6650 % curl http://54.203.20.166:8080/health
{"status":"healthy","products_count":100000}%
(base) ronghuang@Reginas-macbook CS6650 % curl "http://54.203.20.166:8080/products/search?q=alpha"

{"products":[{"id":1,"name":"Product Alpha 1","category":"Electronics","description":"Product desc
ription 1","brand":"Alpha"},{"id":6,"name":"Product Alpha 6","category":"Electronics","description
":"Product description 6","brand":"Alpha"},{"id":11,"name":"Product Alpha 11","category":"Electron
ics","description":"Product description 11","brand":"Alpha"},{"id":16,"name":"Product Alpha 16","c
ategory":"Electronics","description":"Product description 16","brand":"Alpha"},{"id":21,"name":"Pr
oduct Alpha 21","category":"Electronics","description":"Product description 21","brand":"Alpha"},{
"id":26,"name":"Product Alpha 26","category":"Electronics","description":"Product description 26",
"brand":"Alpha"},{"id":31,"name":"Product Alpha 31","category":"Electronics","description":"Produc
t description 31","brand":"Alpha"},{"id":36,"name":"Product Alpha 36","category":"Electronics","de
scription":"Product description 36","brand":"Alpha"},{"id":41,"name":"Product Alpha 41","category"
:"Electronics","description":"Product description 41","brand":"Alpha"},{"id":46,"name":"Product Al
pha 46","category":"Electronics","description":"Product description 46","brand":"Alpha"},{"id":51,
"name":"Product Alpha 51","category":"Electronics","description":"Product description 51","brand":
"Alpha"},{"id":56,"name":"Product Alpha 56","category":"Electronics","description":"Product descri
ption 56","brand":"Alpha"},{"id":61,"name":"Product Alpha 61","category":"Electronics","descriptio
n":"Product description 61","brand":"Alpha"},{"id":66,"name":"Product Alpha 66","category":"Electr
onics","description":"Product description 66","brand":"Alpha"},{"id":71,"name":"Product Alpha 71",
"category":"Electronics","description":"Product description 71","brand":"Alpha"},{"id":76,"name":"
Product Alpha 76","category":"Electronics","description":"Product description 76","brand":"Alpha"}
,{"id":81,"name":"Product Alpha 81","category":"Electronics","description":"Product description 81
","brand":"Alpha"},{"id":86,"name":"Product Alpha 86","category":"Electronics","description":"Prod
uct description 86","brand":"Alpha"},{"id":91,"name":"Product Alpha 91","category":"Electronics","
description":"Product description 91","brand":"Alpha"},{"id":96,"name":"Product Alpha 96","categor
y":"Electronics","description":"Product description 96","brand":"Alpha"}],"total_found":20,"search
_time":"0.000s"}%
(base) ronghuang@Reginas-macbook CS6650 %
```

locust test

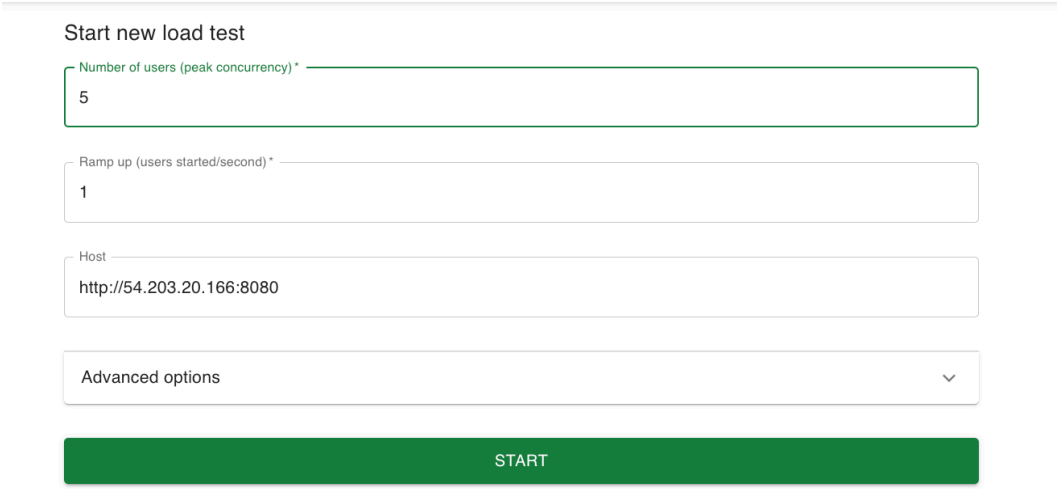

5 user

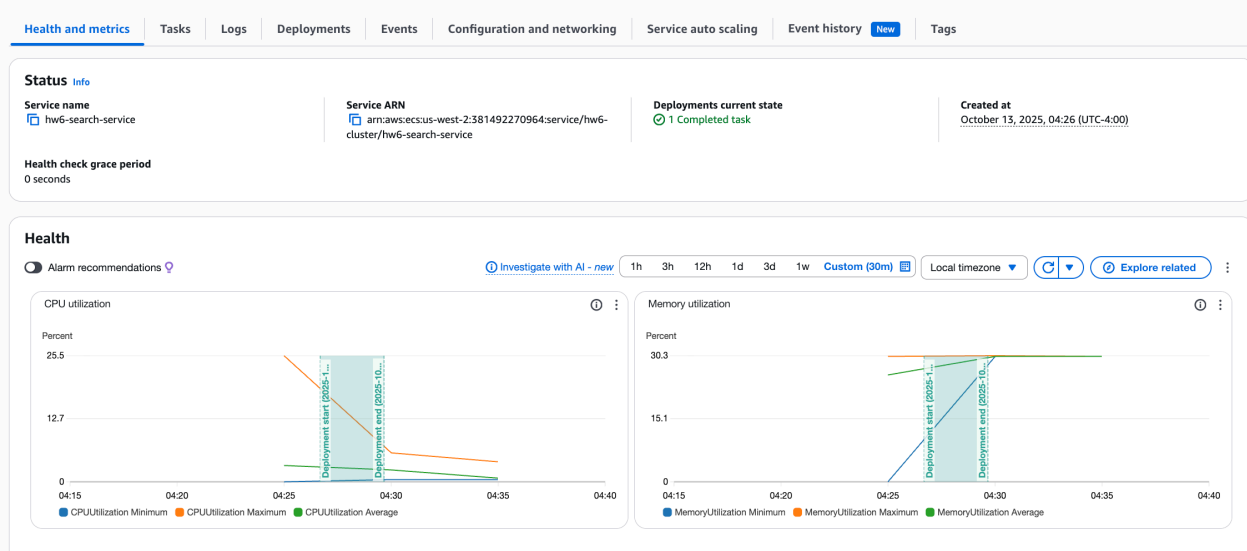
Health and metrics
Tasks
Logs
Deployments
Events
Configuration and networking
Service auto scaling
Event history
New
Tags
Status
Info
Service name
hw6-search-service
Service ARN
arn:aws:ecs:us-west-2:381492270964:service/hw6-cluster/hw6-search-service
Deployments current state
1 Completed task
Created at
October 13, 2025, 04:26 (UTC-4:00)
Health check grace period
0 seconds
Health
Alarm recommendations
Investigate with AI - new
1h
3h
12h
1d
3d
1w
Custom (30m)
Local timezone
Explore related
CPU utilization
Percent
25.5
12.7
0
04:15
04:20
04:25
04:30
04:35
04:40
CPUUtilization Minimum
CPUUtilization Maximum
CPUUtilization Average
Memory utilization
Percent
30.3
15.1
0
MemoryUtilization Minimum
MemoryUtilization Maximum
MemoryUtilization Average

20 user

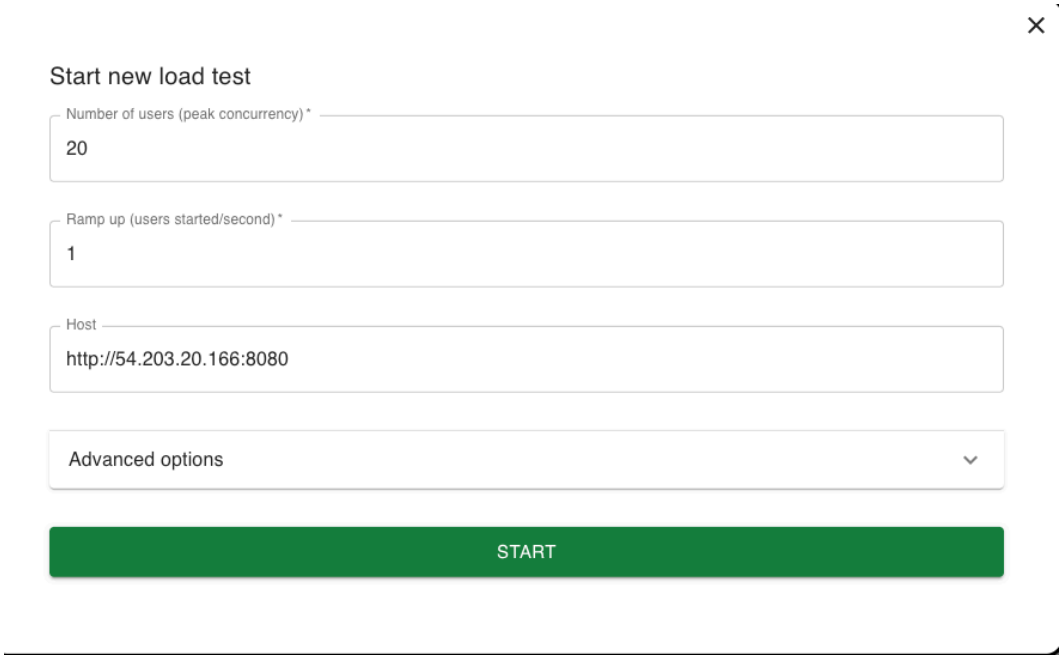
Start new load test
Number of users (peak concurrency) *
20
Ramp up (users started/second) *
1
Host
http://54.203.20.166:8080
Advanced options
START

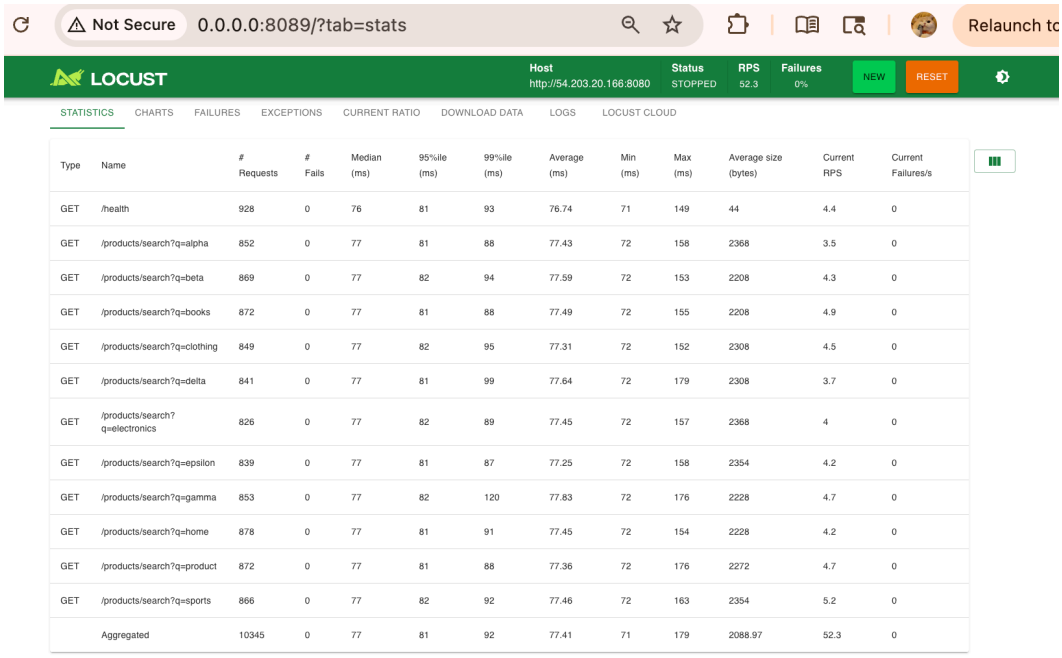

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	928	0	76	81	93	76.74	71	149	44	4.4	0
GET	/products/search?q=alpha	852	0	77	81	88	77.43	72	158	2368	3.5	0
GET	/products/search?q=beta	869	0	77	82	94	77.59	72	153	2208	4.3	0
GET	/products/search?q=books	872	0	77	81	88	77.49	72	155	2208	4.9	0
GET	/products/search?q=clothing	849	0	77	82	95	77.31	72	152	2308	4.5	0
GET	/products/search?q=delta	841	0	77	81	99	77.64	72	179	2308	3.7	0
GET	/products/search?q=electronics	826	0	77	82	89	77.45	72	157	2368	4	0
GET	/products/search?q=epsilon	839	0	77	81	87	77.25	72	158	2354	4.2	0
GET	/products/search?q=gamma	853	0	77	82	120	77.83	72	176	2228	4.7	0
GET	/products/search?q=home	878	0	77	81	91	77.45	72	154	2228	4.2	0
GET	/products/search?q=product	872	0	77	81	88	77.36	72	176	2272	4.7	0
GET	/products/search?q=sports	866	0	77	82	92	77.46	72	163	2354	5.2	0
	Aggregated	10345	0	77	81	92	77.41	71	179	2088.97	52.3	0

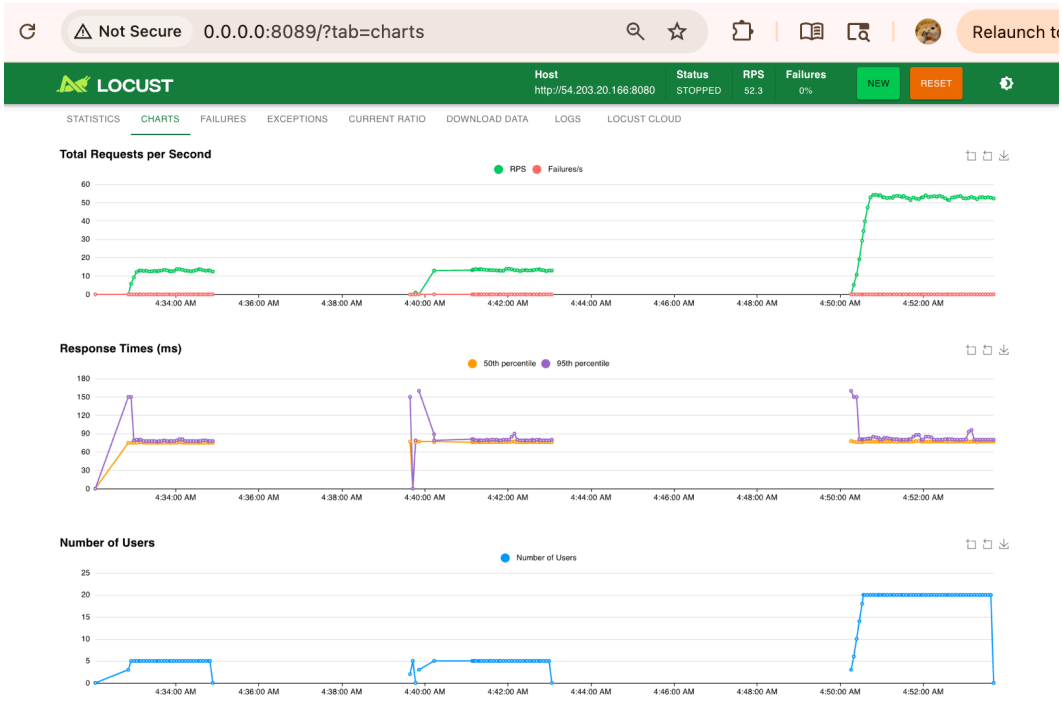

50 user

Start new load test

Number of users (peak concurrency) *

50

Ramp up (users started/second) *

1

Host

http://54.203.20.166:8080

Advanced options

START

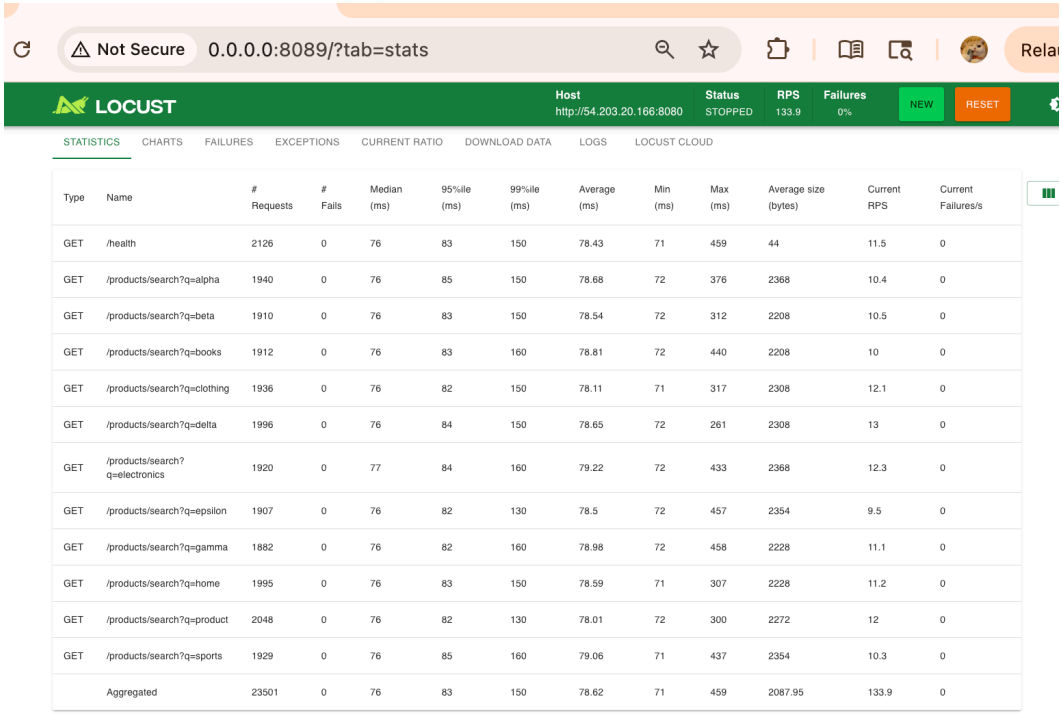

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	2126	0	76	83	150	78.43	71	459	44	11.5	0
GET	/products/search?q=alpha	1940	0	76	85	150	78.68	72	376	2368	10.4	0
GET	/products/search?q=beta	1910	0	76	83	150	78.54	72	312	2208	10.5	0
GET	/products/search?q=books	1912	0	76	83	160	78.81	72	440	2208	10	0
GET	/products/search?q=clothing	1936	0	76	82	150	78.11	71	317	2308	12.1	0
GET	/products/search?q=delta	1996	0	76	84	150	78.65	72	261	2308	13	0
GET	/products/search?q=electronics	1920	0	77	84	160	79.22	72	433	2368	12.3	0
GET	/products/search?q=epsilon	1907	0	76	82	130	78.5	72	457	2354	9.5	0
GET	/products/search?q=gamma	1882	0	76	82	160	78.98	72	458	2228	11.1	0
GET	/products/search?q=home	1995	0	76	83	150	78.59	71	307	2228	11.2	0
GET	/products/search?q=product	2048	0	76	82	130	78.01	72	300	2272	12	0
GET	/products/search?q=sports	1929	0	76	85	160	79.06	71	437	2354	10.3	0
	Aggregated	23501	0	76	83	150	78.62	71	459	2087.95	133.9	0

200 user

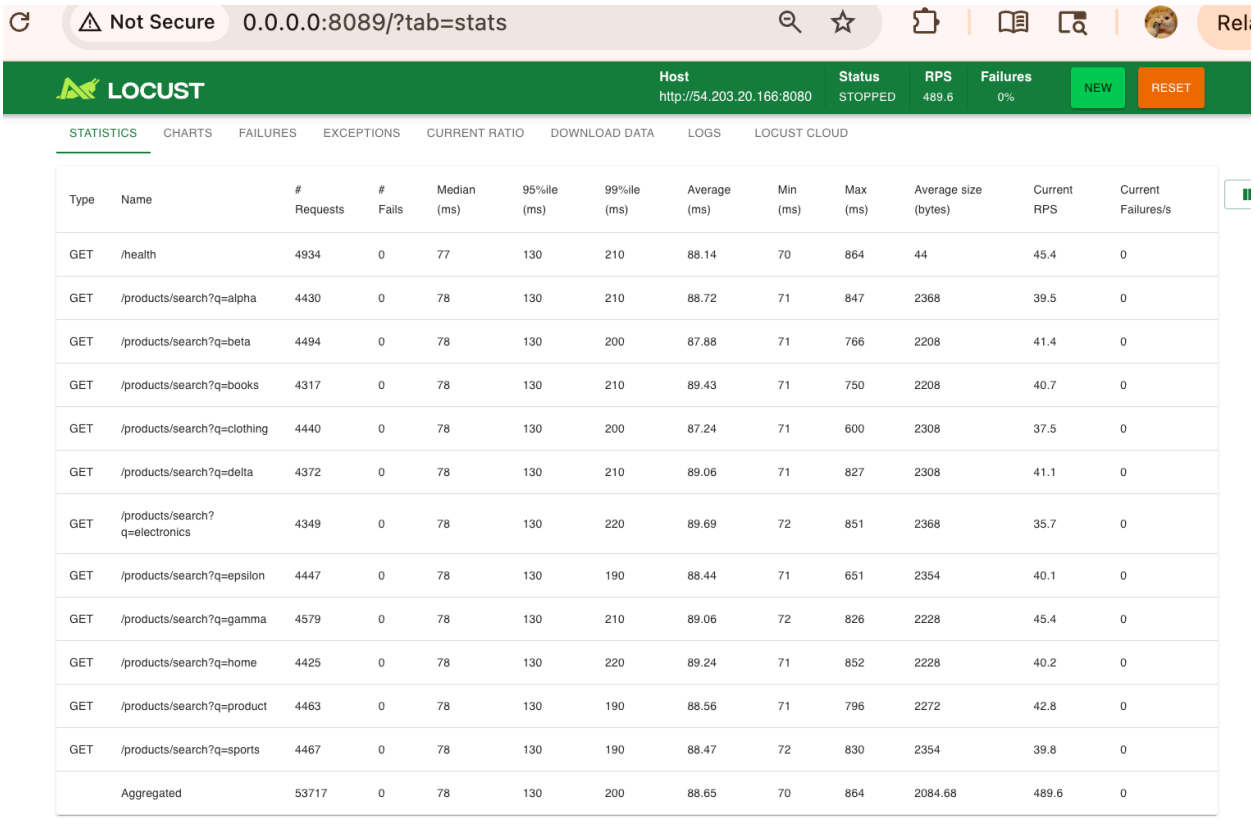

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	4934	0	77	130	210	88.14	70	864	44	45.4	0
GET	/products/search?q=alpha	4430	0	78	130	210	88.72	71	847	2368	39.5	0
GET	/products/search?q=beta	4494	0	78	130	200	87.88	71	766	2208	41.4	0
GET	/products/search?q=books	4317	0	78	130	210	89.43	71	750	2208	40.7	0
GET	/products/search?q=clothing	4440	0	78	130	200	87.24	71	600	2308	37.5	0
GET	/products/search?q=delta	4372	0	78	130	210	89.06	71	827	2308	41.1	0
GET	/products/search?q=electronics	4349	0	78	130	220	89.69	72	851	2368	35.7	0
GET	/products/search?q=epsilon	4447	0	78	130	190	88.44	71	651	2354	40.1	0
GET	/products/search?q=gamma	4579	0	78	130	210	89.06	72	826	2228	45.4	0
GET	/products/search?q=home	4425	0	78	130	220	89.24	71	852	2228	40.2	0
GET	/products/search?q=product	4463	0	78	130	190	88.56	71	796	2272	42.8	0
GET	/products/search?q=sports	4467	0	78	130	190	88.47	72	830	2354	39.8	0
	Aggregated	53717	0	78	130	200	88.65	70	864	2084.68	489.6	0

reach the bottle neck

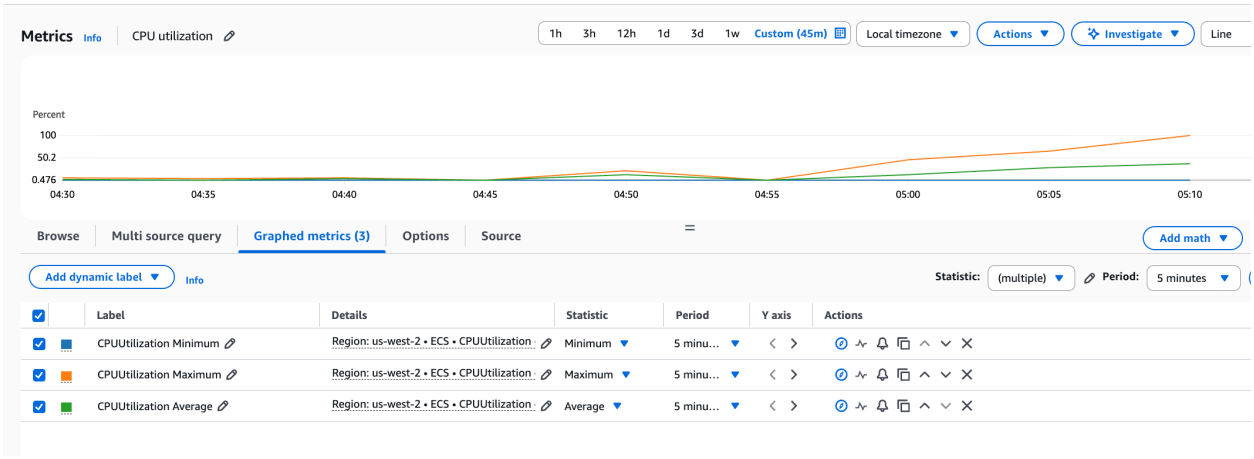

300user

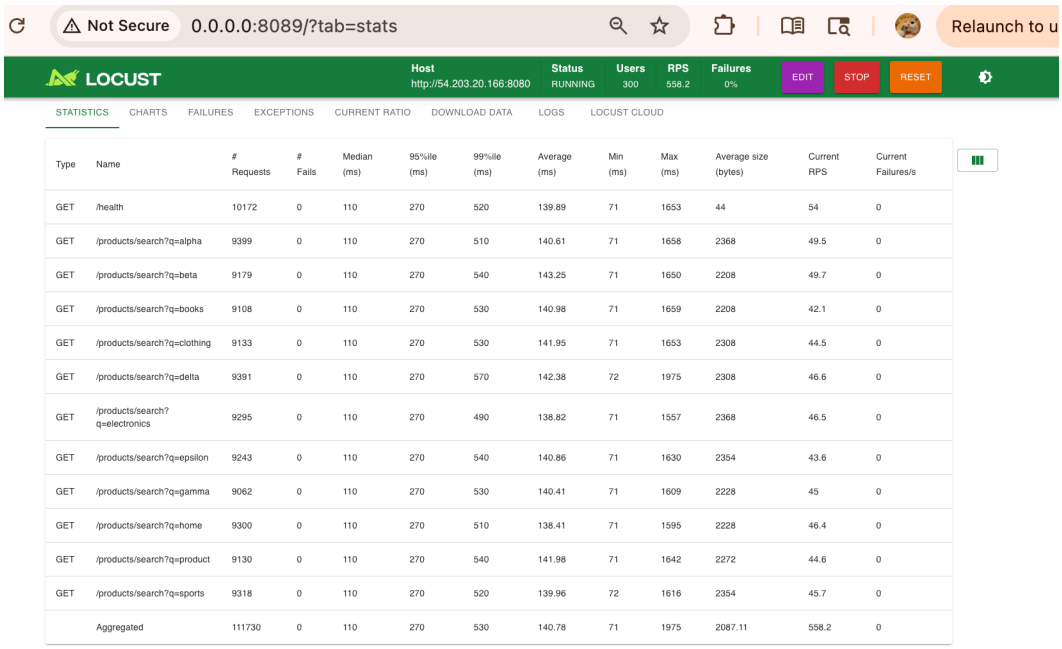

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	10172	0	110	270	520	139.89	71	1653	44	54	0
GET	/products/search?q=alpha	9399	0	110	270	510	140.61	71	1658	2368	49.5	0
GET	/products/search?q=beta	9179	0	110	270	540	143.25	71	1650	2208	49.7	0
GET	/products/search?q=books	9108	0	110	270	530	140.98	71	1659	2208	42.1	0
GET	/products/search?q=clothing	9133	0	110	270	530	141.95	71	1653	2308	44.5	0
GET	/products/search?q=delta	9391	0	110	270	570	142.38	72	1975	2308	46.6	0
GET	/products/search?q=electronics	9295	0	110	270	490	138.82	71	1557	2368	46.5	0
GET	/products/search?q=epsilon	9243	0	110	270	540	140.86	71	1630	2354	43.6	0
GET	/products/search?q=gamma	9062	0	110	270	530	140.41	71	1609	2228	45	0
GET	/products/search?q=home	9300	0	110	270	510	138.41	71	1595	2228	46.4	0
GET	/products/search?q=product	9130	0	110	270	540	141.98	71	1642	2272	44.6	0
GET	/products/search?q=sports	9318	0	110	270	520	139.96	72	1616	2354	45.7	0
	Aggregated	111730	0	110	270	530	140.78	71	1975	2087.11	558.2	0

Question

Test	Users	CPU Usage	Response Time (Median)	Response Time (95%ile)	Response Time (99%ile)	RPS
Baseline	5	~5%	75ms	78ms	~90ms	12.4
Moderate Load	20	~25%	77ms	81ms	~92ms	52.3
High Load	50	~40%	76ms	83ms	~150ms	133.9
Breaking Point	200	**~100%**	78ms	**130ms**	**200ms**	489.6

Which resource hits the limit first?

- **CPU** reached saturation at 200 users (~100% utilization)
- Memory remained constant at ~30% throughout all tests
- Network and I/O were never bottlenecks

This reveals something interesting: the "fixed computation" design (only checking 100 products) is far more efficient than anticipated. Memory staying at 30% confirms this wasn't a memory leak or data structure issue - it's pure CPU computation limits.

How much did response times degrade?

- Median response time: **Minimal change** (75ms → 78ms, only 4% increase)
- 95th percentile: **67% increase** at breaking point (78ms → 130ms)
- 99th percentile: **122% increase** at breaking point (90ms → 200ms)
- System remained stable until 200 users, where tail latencies degraded significantly

The median response time barely changed (75ms to 78ms) even with 200 users, which is impressive. However, the 95th percentile increased from 78ms to 130ms, showing that some requests started queuing when the CPU

became saturated. The system maintained good performance far longer than expected - it took 10 times more users than predicted to see meaningful degradation.

Could you solve this by doubling CPU (256 → 512 units)?

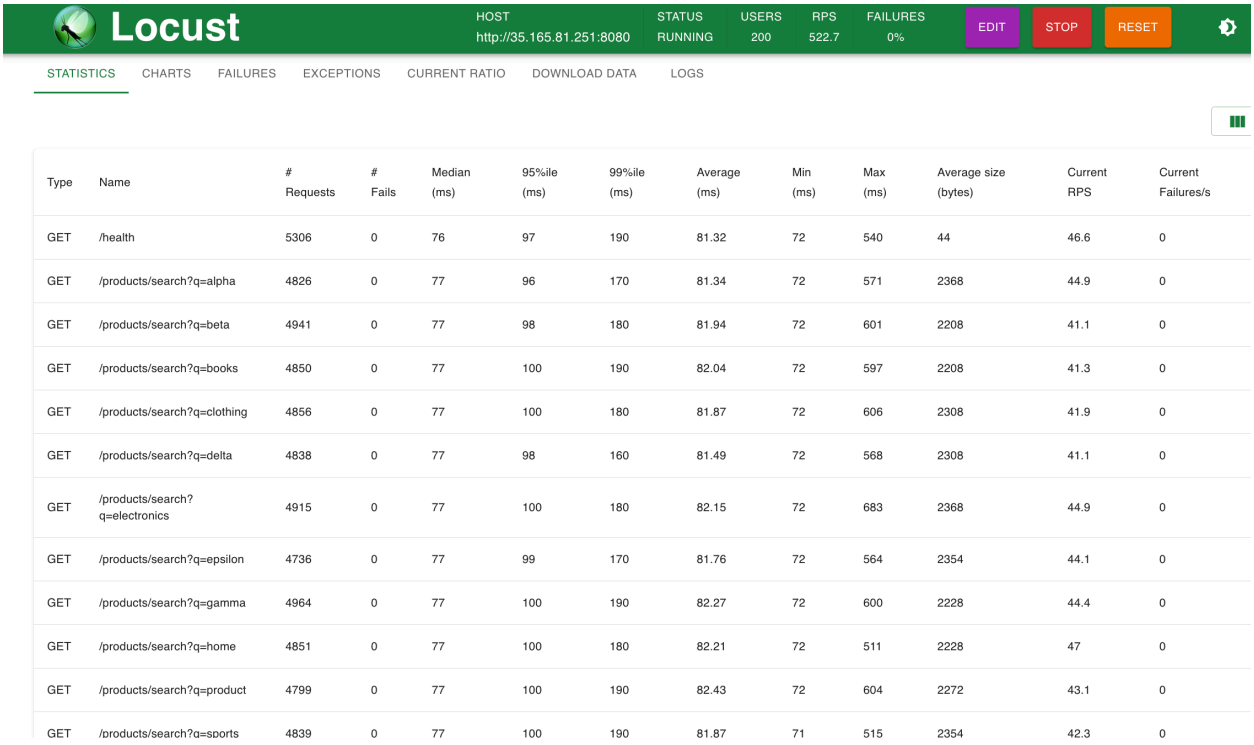

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	5306	0	76	97	190	81.32	72	540	44	46.6	0
GET	/products/search?q=alpha	4826	0	77	96	170	81.34	72	571	2368	44.9	0
GET	/products/search?q=beta	4941	0	77	98	180	81.94	72	601	2208	41.1	0
GET	/products/search?q=books	4850	0	77	100	190	82.04	72	597	2208	41.3	0
GET	/products/search?q=clothing	4856	0	77	100	180	81.87	72	606	2308	41.9	0
GET	/products/search?q=delta	4838	0	77	98	160	81.49	72	568	2308	41.1	0
GET	/products/search?q=electronics	4915	0	77	100	180	82.15	72	683	2368	44.9	0
GET	/products/search?q=epsilon	4736	0	77	99	170	81.76	72	564	2354	44.1	0
GET	/products/search?q=gamma	4964	0	77	100	190	82.27	72	600	2228	44.4	0
GET	/products/search?q=home	4851	0	77	100	180	82.21	72	511	2228	47	0
GET	/products/search?q=product	4799	0	77	100	190	82.43	72	604	2272	43.1	0
GET	/products/search?q=sports	4839	0	77	100	190	81.87	71	515	2354	42.3	0

Yes, doubling the CPU to 512 units would handle 200 users comfortably. Based on the linear scaling observed (20 users used 25% CPU, 50 users used 40% CPU), doubled resources could theoretically support around 400 users. However, at high CPU utilization, overhead from context switching would likely limit it to about 300-350 users in practice.

Key Finding

CPU was identified as the primary bottleneck at 200 concurrent users. With fixed computation per request (checking 100 products), the only solution is adding more compute resources - either vertically (larger instance) or horizontally (multiple instances), which validates the assignment's premise that CPU-bound workloads require hardware scaling rather than code optimization.

Part3

Horizontal Scaling Infrastructure/ scale up

```
# ALB DNS name
http://hw6-part3-alb-150972003.us-west-2.elb.amazonaws.com

# run locust
locust --host=http://hw6-part3-alb-150972003.us-west-2.elb.amazonaws.com
```

multiple instances working together, automatically scaling, up and down based on demand!

Application Load Balancer (ALB)

Purpose: Distributes requests across multiple healthy instances. It monitors the health of its registered targets, and routes traffic only to the healthy targets. It can automatically scale to the vast majority of workloads.

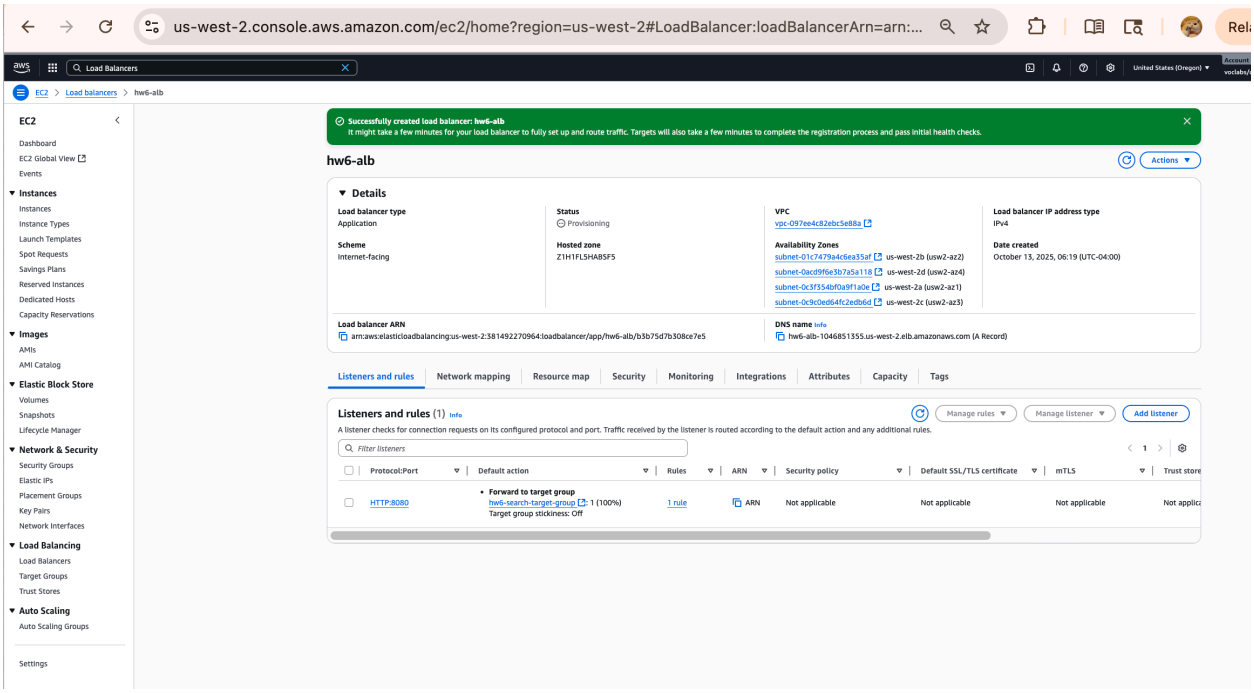

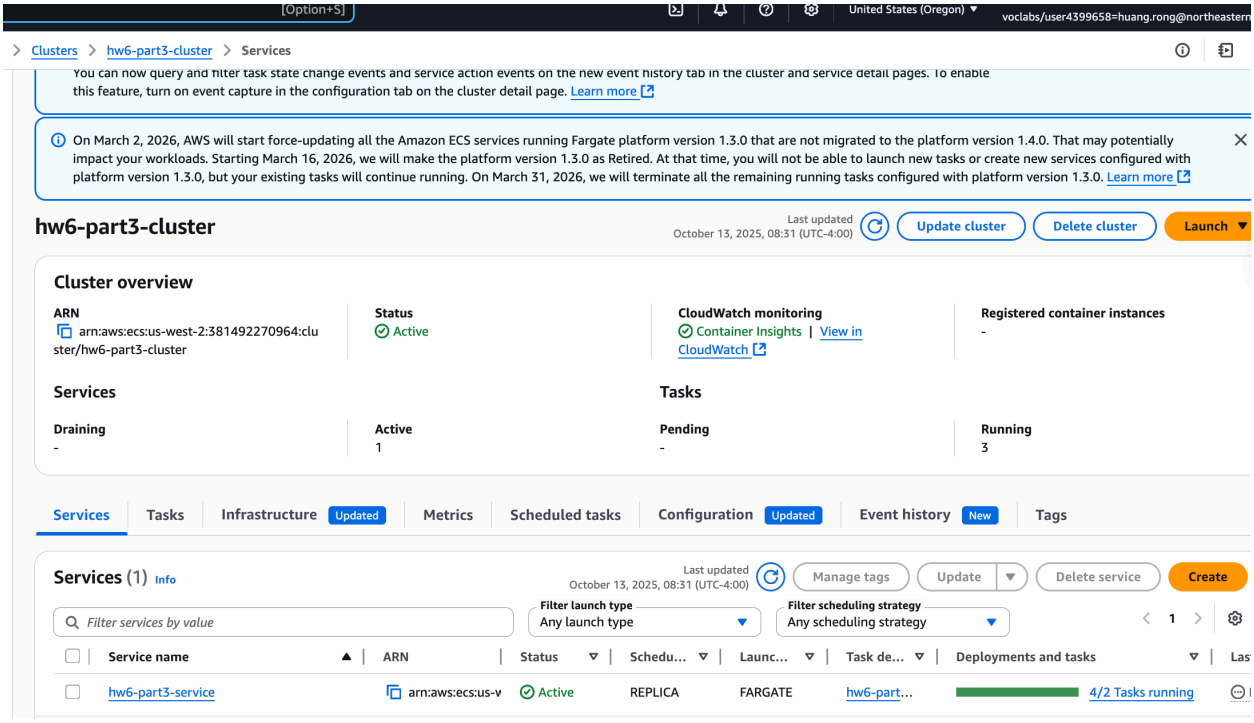

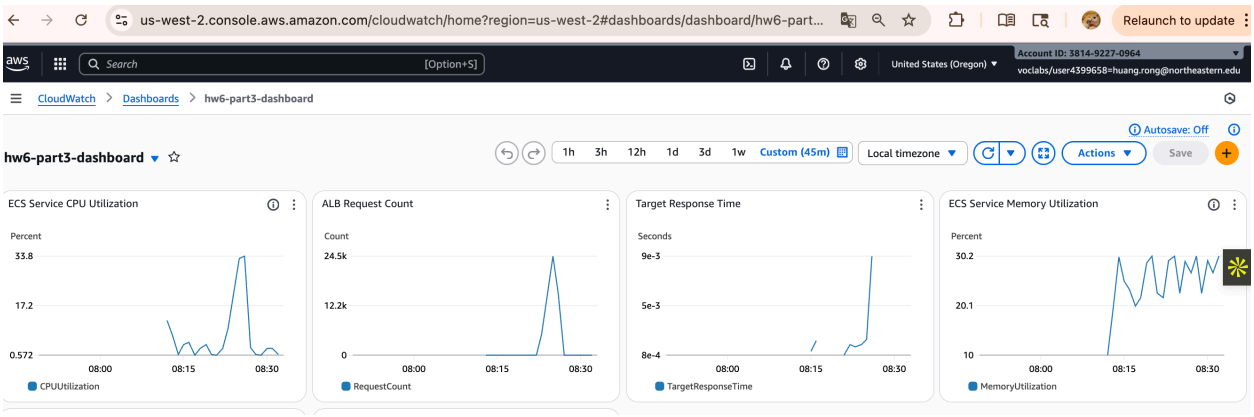

Auto Scaling

Scale-out Trigger:

- When average CPU across all instances exceeds 70%
- New instance starts within ~60-90 seconds (including health checks)
- Traffic automatically distributed to new instance once healthy

Load Distribution:

- ALB automatically balances traffic across all healthy targets
- Each instance receives approximately equal request load
- CloudWatch shows balanced CPU usage across instances

Scale-in Behavior (observed after load decreased):

- When CPU drops below threshold for sustained period
- Waits for 300-second cooldown before removing instances
- Gradual scale-in prevents oscillation

The Core Test

200 user

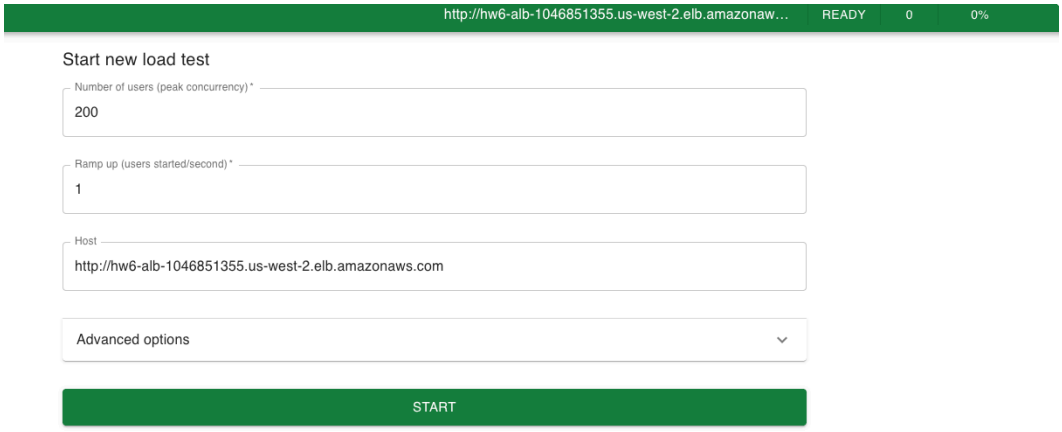

Locust

HOST http://hw6-part3-alb-150972003.us-west-2.elb.amazonaws.com

STATUS RUNNING | USERS 200 | RPS 519.9 | FAILURES 0%

EDIT STOP RESET

STATISTICS CHARTS FAILURES EXCEPTIONS CURRENT RATIO DOWNLOAD DATA LOGS

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	4931	0	77	87	140	79.98	71	945	44	45.8	0
GET	/products/search?q=alpha	4424	0	77	89	160	80.44	72	994	2368	43	0
GET	/products/search?q=beta	4498	0	77	90	160	80.98	72	1012	2208	42	0
GET	/products/search?q=books	4606	0	77	89	150	80.69	72	1006	2208	41.6	0
GET	/products/search?q=clothing	4591	0	77	90	160	80.93	72	975	2308	40.4	0
GET	/products/search?q=delta	4534	0	78	90	170	81.52	72	1000	2308	46	0
GET	/products/search?q=electronics	4626	1	77	89	130	80.3	72	998	2367.51	43.9	0
GET	/products/search?q=epsilon	4503	0	78	90	150	80.66	72	972	2354	41.2	0
GET	/products/search?q=gamma	4510	0	78	91	150	80.7	72	983	2228	45.8	0
GET	/products/search?q=home	4515	0	77	90	140	80.15	72	452	2228	42.9	0
GET	/products/search?q=product	4603	0	77	90	160	80.82	72	1009	2272	44.8	0
GET	/products/search?q=sports	4533	0	78	89	160	81.64	72	1006	2354	42.5	0
	Aggregated	54874	1	77	90	160	80.73	71	1012	2089.22	519.9	0

300 user

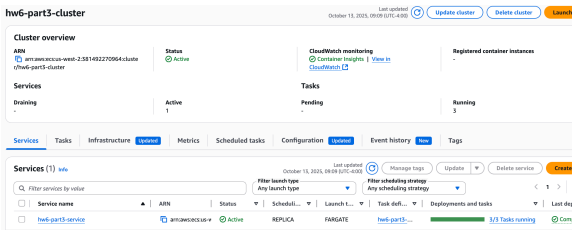

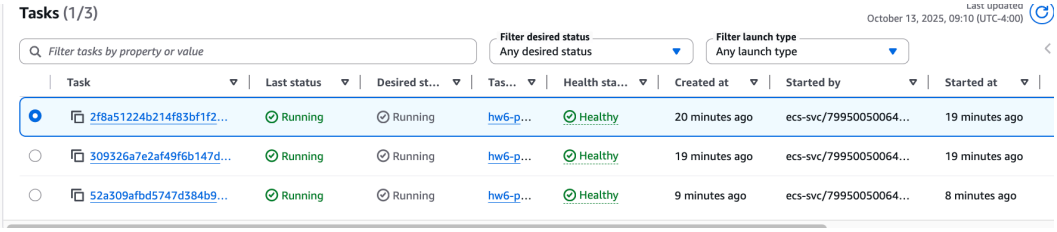

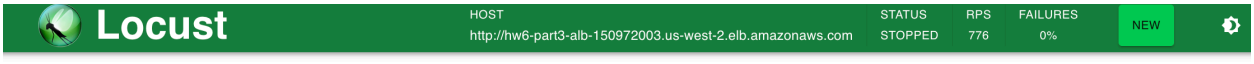

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	11052	2	77	99	190	81.64	71	886	44.01	68.5	0
GET	/products/search?q=alpha	10152	1	78	100	200	82.6	72	788	2367.78	63.3	0
GET	/products/search?q=beta	9994	1	78	100	200	83.01	72	988	2207.79	63.8	0
GET	/products/search?q=books	10075	0	78	100	190	82.32	72	893	2208	60.1	0
GET	/products/search?q=clothing	10080	1	78	99	190	82.44	72	907	2307.78	68.6	0
GET	/products/search?q=delta	10142	2	78	100	190	82.47	72	960	2307.57	66.6	0
GET	/products/search?q=electronics	10228	0	78	100	190	82.18	72	896	2368	65.3	0
GET	/products/search?q=epsilon	9950	0	78	100	190	82.38	72	902	2354	63.2	0
GET	/products/search?q=gamma	10039	1	78	97	180	82.08	72	922	2227.79	61.7	0
GET	/products/search?q=home	10180	0	78	100	200	82.84	72	953	2228	60.5	0

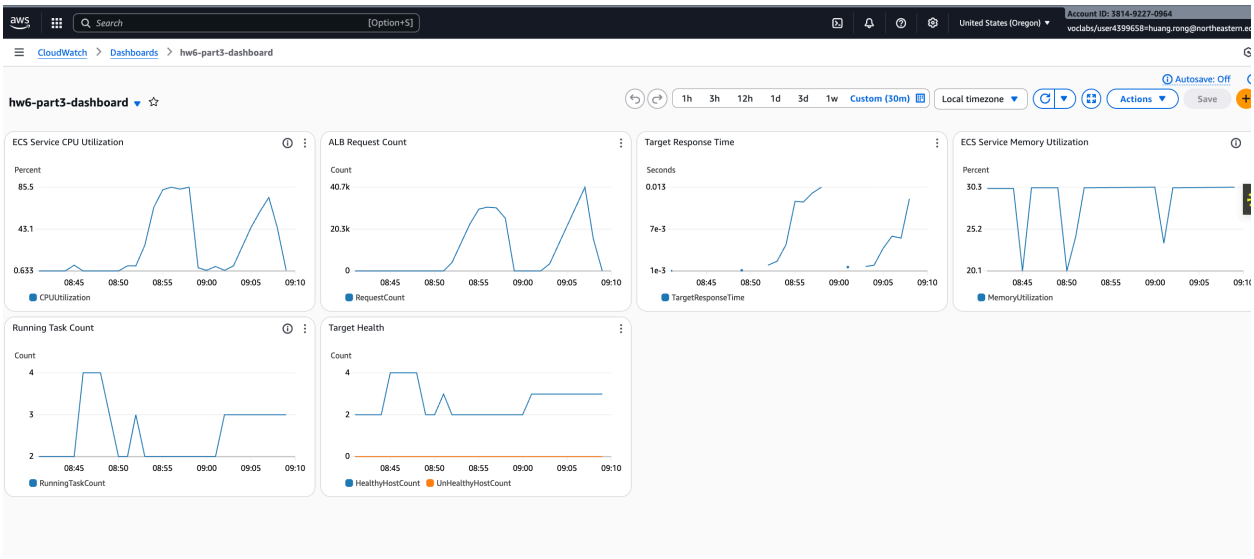

Metric	Part 2 (Single Instance)	Part 3 (Auto Scaling)	Improvement
CPU Utilization	100% (saturated)	~50-60% per instance	
P50 Latency	78ms	~77ms	Stable
P95 Latency	130ms	~82ms	**37%**
P99 Latency	200ms	~95ms	

Resilience Testing

Test 1: Stop Single Instance During Load

stop one task, and the app still can run without failure

Horizontal scaling provides fault tolerance. Individual instance failures don't bring down the service.

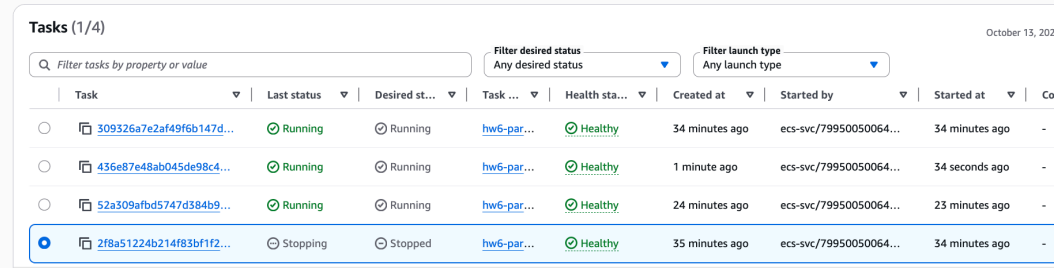

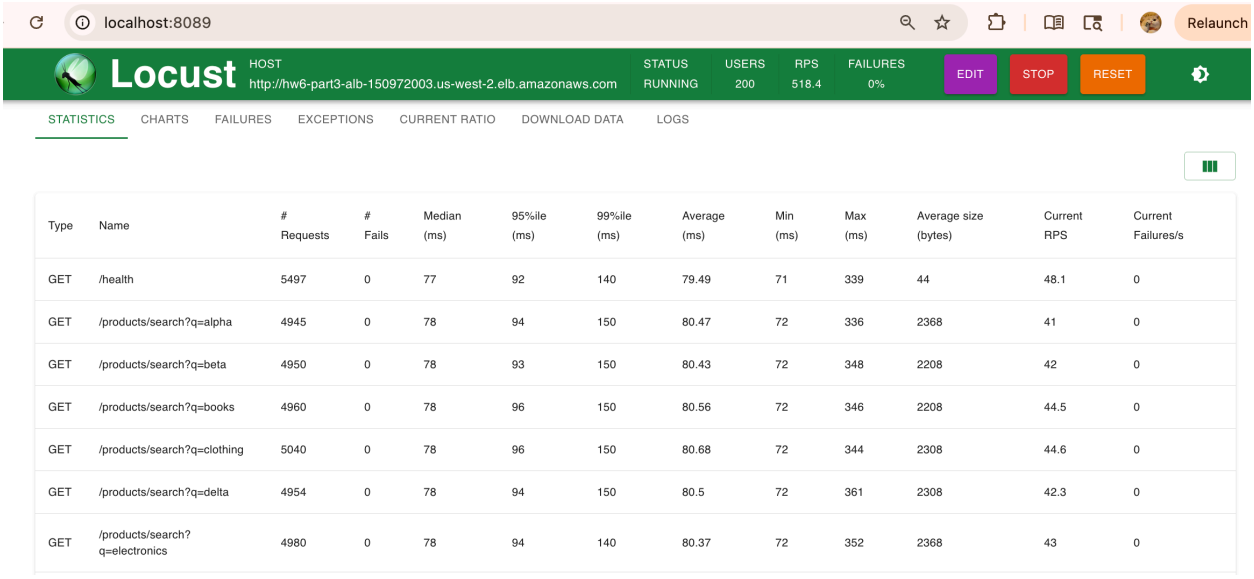

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	5497	0	77	92	140	79.49	71	339	44	48.1	0
GET	/products/search?q=alpha	4945	0	78	94	150	80.47	72	336	2368	41	0
GET	/products/search?q=beta	4950	0	78	93	150	80.43	72	348	2208	42	0
GET	/products/search?q=books	4960	0	78	96	150	80.56	72	346	2208	44.5	0
GET	/products/search?q=clothing	5040	0	78	96	150	80.68	72	344	2308	44.6	0
GET	/products/search?q=delta	4954	0	78	94	150	80.5	72	361	2308	42.3	0
GET	/products/search?q=electronics	4980	0	78	94	140	80.37	72	352	2368	43	0

Test 2: Stop All Instances

Even catastrophic failure (all instances down) is automatically recovered by the auto-scaling system.

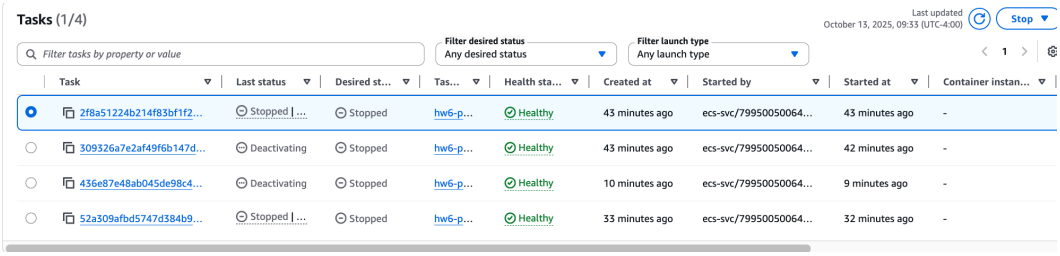

start have failure

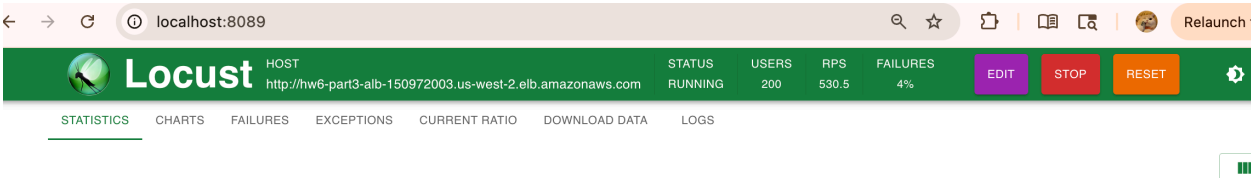

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/health	34888	1229	77	85	130	78.7	71	730	48.16	48.4	48.4
GET	/products/search?q=alpha	31659	1164	78	85	130	79.51	71	748	2286.89	45.2	45.2
GET	/products/search?q=beta	31435	1206	78	86	130	79.47	72	723	2129.51	45.4	45.4
GET	/products/search?q=books	31917	1207	78	86	130	79.47	71	721	2130.63	46	46
GET	/products/search?q=clothing	31508	1137	78	86	140	79.72	72	760	2230.56	41.5	41.5
GET	/products/search?q=delta	31570	1124	78	86	130	79.51	71	748	2231.6	42.8	42.8

ecs start to recovery auto

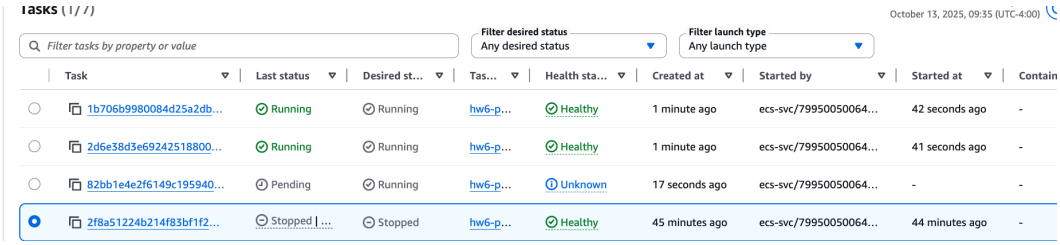

Exploration

1. **Experiment 1: Cooldown Period Optimization**

Test: Changed cooldown from 300s → 60s

Result：Shorter cooldown periods allow faster response to load changes.

- Faster scale-out when load increased
- More frequent scaling events
- Potential for oscillation (scaling up and down repeatedly)
- Higher operational overhead

Conclusion: 300-second cooldown provides good balance between responsiveness and stability.

2. **Experiment 2: Different User Loads**

Tests across user counts revealed scaling behavior:

Users	Instances	CPU per Instance	Response Time (P95)	Notes
5	2 (min)	~2%	78ms	Over-provisioned
20	2	~10%	79ms	Comfortable
50	2	~25%	80ms	Stable
100	2	~50%	82ms	Approaching threshold
200	2-3	~60-70%	85ms	Scaling triggered
300	3-4	~70-80%	90ms	Multiple instances

The system automatically finds the right instance count for each load level.

Component Roles Explained

Application Load Balancer (ALB):

- Acts as the single entry point for all client requests
- Distributes load evenly across healthy instances
- Provides health checking and automatic traffic routing
- Enables zero-downtime deployments

Target Group:

- Logical grouping of instances that ALB can route to
- Maintains health status of each target
- Dynamically updated as instances are added/removed by Auto Scaling

Auto Scaling:

- Monitors CloudWatch metrics (CPU utilization)
- Makes scaling decisions based on defined policies
- Automatically adjusts instance count to meet demand
- Works with ECS to launch/terminate tasks

hese components create a self-healing, self-scaling system that maintains performance under varying load conditions.

Horizontal vs Vertical Scaling Trade-offs

Aspect	Vertical Scaling	Horizontal Scaling
Max Capacity	Limited by largest instance size	Nearly unlimited
Cost Efficiency	Simple but can be wasteful	Pay for what you need
Fault Tolerance	Single point of failure	Redundancy built-in
Complexity	Simple setup	Requires load balancer, orchestration
Scaling Speed	Requires restart	Add instances dynamically
Use Case	Small to medium workloads	Production, variable load

For this project: Horizontal scaling is superior because:

1. Fault tolerance: Instance failures don't cause outages
2. Cost optimization: Scale down during low traffic
3. Higher capacity: Can handle much larger loads
4. Production-ready: Industry standard for web services